

Ship Detection in Satellite Imagery

Sam Markelon and Tyler Tucker

Ship Detection

- Airbus Ship Detection Challenge
 - <https://www.kaggle.com/c/airbus-ship-detection/overview>
- Challenge to detect ships in satellite imagery to prevent piracy, monitor trading volume, estimate environmental impact, and more
- Binary classification problem
 - Ship present or ship not present
- Rich and large data set
 - 192556 unique labeled examples
 - 0.78 are “no ship” class, 0.22 are “ship” class

Approach




- Use a convolutional neural network (CNN) to classify the data set
 - CNN
 - CNN with dropout
 - Transfer learning on pretrained image classification model
 - Transfer learning with fine tunings
- Train on HiPerGator
 - Time constraints made it prohibitive to train on entire data set
 - Randomly sample 10000 examples for training set and 2000 for validation set
 - Representative of the entire data set
- Compare the approaches

Data Processing and Feature Engineering

- Process the labeled examples CSV to get format be a listing of the form
`ImageID:classLabel`
 - `process_data.py`
- Sample and split the data set
- Rescale the data by $1. / 255$
- Use the Keras `ImageDataGenerator` class to create data set generators
 - Resize the image to 150x150
 - `x = ImageId, y = class`
 - batch size = 20
 - Binary classification mode

	ImageId	ShipPresent
0	00003e153.jpg	0
1	0001124c7.jpg	0
2	000155de5.jpg	1
3	000194a2d.jpg	1
4	0001b1832.jpg	0
...
192551	fffedbb6b.jpg	0
192552	ffff2aa57.jpg	0
192553	ffff6e525.jpg	0
192554	ffffc50b4.jpg	0
192555	ffffe97f3.jpg	0

192556 rows × 2 columns



CNN Model

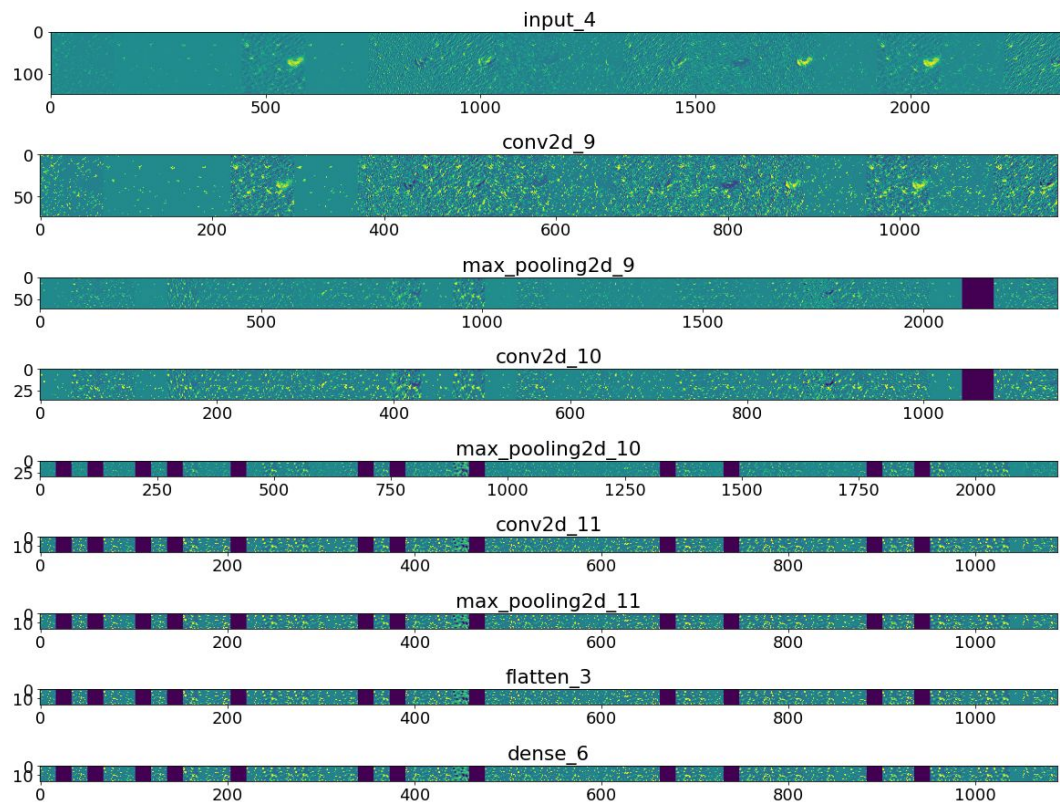
- Three convolution layers using relu activation
- Output layer used sigmoid activation function to encode probability of ship present as scalar in range [0,1].
- Train the model using binary cross-entropy as a loss function with a learning rate of 0.001 over 15 epochs with a batch size of 20.
- Steps per epoch determined by:

$\text{steps} = \text{ceil}(\text{data set size} / \text{batch size})$

Model: "functional_7"


Layer (type)	Output Shape	Param #
input_4 (InputLayer)	[(None, 150, 150, 3)]	0
conv2d_9 (Conv2D)	(None, 148, 148, 16)	448
max_pooling2d_9 (MaxPooling2)	(None, 74, 74, 16)	0
conv2d_10 (Conv2D)	(None, 72, 72, 32)	4640
max_pooling2d_10 (MaxPooling)	(None, 36, 36, 32)	0
conv2d_11 (Conv2D)	(None, 34, 34, 64)	18496
max_pooling2d_11 (MaxPooling)	(None, 17, 17, 64)	0
flatten_3 (Flatten)	(None, 18496)	0
dense_6 (Dense)	(None, 512)	9470464
dense_7 (Dense)	(None, 1)	513
Total params: 9,494,561		
Trainable params: 9,494,561		
Non-trainable params: 0		

Intermediate Image Representation by Layer



CNN Results

- Measure model using accuracy metric
- Training accuracy = 0.9687
- Validation accuracy = 0.8955
- Good model, but indicative of overfitting



CNN with Dropout

- Add a dropout layer with dropout rate of 0.5 before the output layer to combat overfitting
- Makes the problem artificially more difficult for the model
- Train the model using the same hyperparameters as before


Model: "functional_19"

Layer (type)	Output Shape	Param #
input_8 (InputLayer)	[(None, 150, 150, 3)]	0
conv2d_21 (Conv2D)	(None, 148, 148, 16)	448
max_pooling2d_21 (MaxPooling)	(None, 74, 74, 16)	0
conv2d_22 (Conv2D)	(None, 72, 72, 32)	4640
max_pooling2d_22 (MaxPooling)	(None, 36, 36, 32)	0
conv2d_23 (Conv2D)	(None, 34, 34, 64)	18496
max_pooling2d_23 (MaxPooling)	(None, 17, 17, 64)	0
flatten_7 (Flatten)	(None, 18496)	0
dense_14 (Dense)	(None, 512)	9470464
dropout_5 (Dropout)	(None, 512)	0
dense_15 (Dense)	(None, 1)	513
Total params: 9,494,561		
Trainable params: 9,494,561		
Non-trainable params: 0		


CNN with Dropout Results

- Measure model using accuracy metric
- Training accuracy = 0.9193
- Validation accuracy = 0.8870
- Model is slightly less good than before, but we have lessened the amount of overfitting
- If we properly tune the dropout rate, we can likely achieve a model with similar accuracy that does not overfit

Training and validation accuracy



Training and validation loss





Transfer Learning

- Apply a model that is good at making predictions on a related but different dataset
- We use Google's Inception-v3 model trained on the ImageNet data set
- Freeze the model up to the mixed7 layer
 - last feature extraction layer
- Add a flatten layer, a dense layer, a drop our layer (dropout rate =0.2), and an output layer on top of the Inception-v3 layers
 - All feature extraction is done by Inception-v3 model
- Train the model with a learning rate of 0.0001 over 5 epochs
- **Achieve a training accuracy of 0.9126 and a validation accuracy of 0.9050**

Transfer Learning with Fine Tunings

- To see if we can achieve better results, we fine tune the weights of the pre-trained model's top layers on our data set
- We use stochastic gradient descent to optimize the weights using the binary cross-entropy loss function
- Trained using a learning rate of 0.00001, a momentum of 0.9, over 10 epochs.
- **Achieve a training accuracy of 0.9361 and a validation accuracy of 0.9080**
- Best model out of all the ones we tested
 - only slightly better than transfer learning without fine tunings
 - more overfit



Conclusion

- Performance on all models was similar
 - Transfer learning with fine tunings was the the the best with over 90% validation accuracy
 - Squeeze out slightly better performance with further hyperparameter tuning
 - Performed on par with the top submissions on Kaggle
 - Slightly different problem there
 - Had to identify pixels in the image where ship is present, while we did not do this due to time constraints
- Further work
 - Identify pixels where ship is present
 - Multiclass classification (small vessel, shipping vessel, etc.)
 - Explainable model